

# Fundamental of ML Final Project: Handwritten Image Classifier

## Library imports

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
import tensorflow.keras as keras
import cv2
from sklearn.preprocessing import MinMaxScaler, OneHotEncoder
from sklearn.model_selection import train_test_split
import time
```

## Unused Functions

```
In [947... def vert_det(im, kernel):
    vert = cv2.Sobel(im, ddepth=cv2.CV_64F, dx=0, dy=1, ksize=kernel)
    return im-vert

def hor_det(im, kernel):
    hor = cv2.Sobel(im, ddepth=cv2.CV_64F, dx=1, dy=0, ksize=kernel)
    return im-hor

def binarize(im):
    threshold = np.max(im)-.6
    im[im < threshold] = 0
    im[im >= threshold] = 1
    return im

def otsu(im):
    im = im.astype("uint8")
    _, r = cv2.threshold(im, 0, 255, cv2.THRESH_BINARY+cv2.THRESH_OTSU)
    return r

def bounding_box_transform(im):
    tt = np.argmax(im, axis=0)
    leftbound = np.amin(np.nonzero(tt))
    rightbound = np.amax(np.nonzero(tt))
    t = np.argmax(im, axis=1)
    upbound = np.amin(np.nonzero(t))
    downbound = np.amax(np.nonzero(t))
    boundarybox = np.float32([[leftbound-10, upbound-10], [rightbound+10, upbound-10],
                             [leftbound-10, downbound+10], [rightbound+10, downbound+10]])
    newbox = np.float32([[0,0], [300,0], [0,300], [300,300]])
    M = cv2.getPerspectiveTransform(boundarybox, newbox)
    dst = cv2.warpPerspective(im, M, (300,300))
    return dst

# takes in a non-inverted image (non min max normalized)
def remove_lines(z):
    z = z.reshape(300,300)
    z = update_type(z)
```

```

#threshold the image
thresh = cv2.threshold(z, 0, 255, cv2.THRESH_BINARY_INV + cv2.THRESH_OTSU)[1]

# Remove horizontal
horizontal_kernel = cv2.getStructuringElement(cv2.MORPH_RECT, (50,1))
detected_lines = cv2.morphologyEx(thresh, cv2.MORPH_OPEN, horizontal_kernel, iterations=1)
cnts = cv2.findContours(detected_lines, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)
cnts = cnts[0] if len(cnts) == 2 else cnts[1]
for c in cnts:
    cv2.drawContours(z, [c], -1, (255,255,255), 2)

# Repair image
repair_kernel = cv2.getStructuringElement(cv2.MORPH_RECT, (1,6))
result = 255 - cv2.morphologyEx(255 - z, cv2.MORPH_CLOSE, repair_kernel, iterations=1)

#threshold the image
thresh = cv2.threshold(result, 0, 255, cv2.THRESH_BINARY_INV + cv2.THRESH_OTSU)[1]

# Remove vertical
vertical_kernel = cv2.getStructuringElement(cv2.MORPH_RECT, (1,50))
detected_lines = cv2.morphologyEx(thresh, cv2.MORPH_OPEN, vertical_kernel, iterations=1)
cnts = cv2.findContours(detected_lines, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)
cnts = cnts[0] if len(cnts) == 2 else cnts[1]
for c in cnts:
    cv2.drawContours(result, [c], -1, (255,255,255), 2)

# Repair image
repair_kernel = cv2.getStructuringElement(cv2.MORPH_RECT, (6,1))
result = 255 - cv2.morphologyEx(255 - result, cv2.MORPH_CLOSE, repair_kernel, iterations=1)

return result

```

## Used Functions

In [948...

```

def show(im):
    plt.imshow(im, cmap='gray')
    plt.show()

def invert(im):
    return (im*-1)+255

def min_max_scale(im):
    im = MinMaxScaler().fit_transform(im.ravel().reshape(-1,1))
    return im.reshape(300,300)

def brighten(im):
    m = np.max(im)-.3
    im[im >= m] += 0.3
    im[im < m] -= 0.3
    return im

def blur(im, kernel):
    im = im.astype("uint8")
    im = cv2.medianBlur(im, kernel)
    return im

def morph_close(im, kernel, i):

```

```

    return cv2.morphologyEx(im, cv2.MORPH_CLOSE, kernel, iterations=i)

def morph_open(im, kernel, i):
    return cv2.morphologyEx(im, cv2.MORPH_OPEN, kernel, iterations=i)

def morph_dilate(im, kernel, i):
    return cv2.dilate(im, kernel, iterations=i)

def morph_erode(im, kernel, i):
    return cv2.erode(im, kernel, iterations=i)

def transform(im):
    dest = np.float32([[0,0],[300,0],[0,300],[300,300]])
    source = np.float32([[25,25],[275,25],[25,275],[275,275]])
    res = cv2.getPerspectiveTransform(source,dest)
    return cv2.warpPerspective(im, res, (300,300))

```

In [991...

```

def preprocess(ims):
    r = []
    for im in ims.T:
        im = im.reshape((300,300))
        im = transform(im)
        im = invert(im)
        im = blur(im, 5)
        im = morph_erode(im, (3,3), 5)
        im = morph_dilate(im, (3,3), 5)
        im = morph_close(im, (3,3), 5)
        im = morph_open(im, (3,3), 5)
        im = min_max_scale(im)
        im = brighten(im)
        im = min_max_scale(im)
        im = cv2.resize(im, (100,100))
        im = np.stack((im,)*3, axis=-1)
        r.append(im)
    return np.array(r)

def augment(ims, labels):
    result = []
    n = []
    o = []
    t = []
    for im in ims:
        im = cv2.rotate(im, cv2.cv2.ROTATE_90_CLOCKWISE)
        n.append(im)
        im = cv2.rotate(im, cv2.cv2.ROTATE_90_CLOCKWISE)
        o.append(im)
        im = cv2.rotate(im, cv2.cv2.ROTATE_90_CLOCKWISE)
        t.append(im)
    result.append(ims)
    result.append(n)
    result.append(o)
    result.append(t)
    labels_new = np.concatenate((labels,labels,labels,labels))
    return np.array(result).reshape(-1,100,100), labels_new

```

## Data imports

```
In [103... data = np.load("data_train.npy")
labels = np.load("labels_train.npy")
```

```
In [103... cnn_input, labels_new = augment(preprocess(data), labels)
```

```
In [107... cnn_input = cnn_input.reshape(-1,100,100,3)
```

```
In [100... # train test split twice to get the splits
# train_x, inter_x, train_y, inter_y = train_test_split(data_rot_augment, labels_augmen
x_train, x_inter, labels_train, labels_inter = train_test_split(cnn_input, labels_new,
```

```
In [100... x_val, x_test, labels_val, labels_test = train_test_split(x_inter, labels_inter, train_
```

Ensuring that data is in the proper format for input into Xception

```
In [100... x_train = x_train.reshape(-1, 100, 100, 3)
x_val = x_val.reshape(-1, 100, 100, 3)
x_test = x_test.reshape(-1, 100, 100, 3)
```

```
In [101... x_train.shape, x_val.shape, x_test.shape
```

```
Out[101... ((18816, 100, 100, 3), (4032, 100, 100, 3), (4032, 100, 100, 3))
```

```
In [101... labels_train_ohe = OneHotEncoder().fit_transform(labels_train.reshape(-1,1)).toarray()
labels_val_ohe = OneHotEncoder().fit_transform(labels_val.reshape(-1,1)).toarray()
labels_test_ohe = OneHotEncoder().fit_transform(labels_test.reshape(-1,1)).toarray()
```

```
In [101... # follow the following format
cnn_dataset = (x_train, labels_train_ohe, x_val, labels_val_ohe, x_test, labels_test_oh
```

## CNN training and architecture

```
In [966... # Let's create a custom callback class
class PerfEvalCustomCallback(keras.callbacks.Callback):

    def __init__(self, perf_data):
        self.perf_data = perf_data

    # we define the on_epoch_end callback and save the loss and accuracy in perf_data
    def on_epoch_end(self, epoch, logs=None):
        self.perf_data[epoch,0] = logs['loss']
        self.perf_data[epoch,1] = logs['accuracy']
        self.perf_data[epoch,2] = logs['val_loss']
        self.perf_data[epoch,3] = logs['val_accuracy']

    def get_perf_data():
        return self.perf_data
```

In [967...

```

# Plot the model's performance during training (across epochs)
def plot_training_perf(train_loss, train_acc, val_loss, val_acc, fs=(8,5)):
    plt.figure(figsize=fs)

    assert train_loss.shape == val_loss.shape and train_loss.shape == val_acc.shape and

    # assume we have one measurement per epoch
    num_epochs = train_loss.shape[0]
    epochs = np.arange(0, num_epochs)

    # Can you figure out why this makes sense? Why remove -0.5?
    plt.plot(epochs-0.5, train_loss, 'm', linewidth=2, label='Loss (Training)')
    plt.plot(epochs-0.5, train_acc, 'r--', linewidth=2, label='Accuracy (Training)')

    plt.plot(epochs, val_loss, 'g', linewidth=2, label='Loss (Validation)')
    plt.plot(epochs, val_acc, 'b:', linewidth=2, label='Accuracy (Validation)')

    plt.xlim([0, num_epochs])
    plt.ylim([0, 1.05])

    plt.legend()

    plt.show()

```

In [101...

```

# Customize this function as you like but makes sure it is implemented correctly.
# Note: If you need to change the method definition to add more arguments, make sure to
# the new arguments optional (and have a sensible default value)

from sklearn.metrics import classification_report, balanced_accuracy_score

def evaluate_model(name, model, eval_data,
                  plot_training=True, evaluate_on_test_set=True):

    # unpack the stuff
    perf_data, dataset = eval_data
    train_x, train_y, val_x, val_y, test_x, test_y = dataset

    # get predictions from the model
    train_preds = model.predict(train_x)
    val_preds = model.predict(val_x)

    # measure the accuracy (as categorical accuracy since we have a softmax layer)
    catacc_metric = keras.metrics.CategoricalAccuracy()
    catacc_metric.update_state(train_y, train_preds)
    train_acc = catacc_metric.result()

    catacc_metric = keras.metrics.CategoricalAccuracy()
    catacc_metric.update_state(val_y, val_preds)
    val_acc = catacc_metric.result()
    print('[{}] Training Accuracy: {:.3f}%, Validation Accuracy: {:.3f}%'.format(name,
    train_acc, val_acc))

    if plot_training:
        plot_training_perf(perf_data[:,0], perf_data[:,1], perf_data[:,2], perf_data[:,3],
                           train_acc, val_acc)

    if evaluate_on_test_set:

```

```

### Evaluate the model on test data and put the results in 'test_loss', 'test_a
###* put your code here (~1-2 lines) *###

test_preds = model.predict(test_x)

accuracy_obj = keras.metrics.CategoricalAccuracy()
accuracy_obj.update_state(test_y, test_preds)
test_acc = accuracy_obj.result()

test_loss, _ = model.evaluate(test_x, test_y) #the test_acc could also be extra

cat_loss_obj = tf.keras.losses.CategoricalCrossentropy()
test_ce_loss = cat_loss_obj(test_y, test_preds)

print('[{}] Test loss: {:.5f}; test accuracy: {:.3f}%'.format(name, test_loss,
print('[{}] Test cross entropy loss: {:.5f}'.format(name, test_ce_loss))

# You can add stuff here if you want
###* put your code here (0+ lines) *###

return

```

```
In [ ]: evaluate_model("CNN", cnn_model, eval_data)
```

```
In [108...
base_model = keras.applications.Xception(
weights='imagenet',
input_shape=(100,100,3),
include_top=False)

base_model.trainable=False

```

```
In [109...
inputs = keras.Input(shape=(100,100,3))
x = base_model.output
x = keras.layers.GlobalAveragePooling2D()(x)
x = keras.layers.Dense(200, activation='relu')(x)
x = keras.layers.Dense(100, activation='relu')(x)
x = keras.layers.Dense(50, activation='relu')(x)
outputs = keras.layers.Dense(10, activation='softmax')(x)
model = keras.Model(inputs=base_model.inputs, outputs=outputs)

```

```
In [109...
for layer in model.layers:
    layer.trainable = True

```

```
In [109...
model.compile(loss='categorical_crossentropy', optimizer='nadam', metrics=['accuracy'])

```

```
In [108...
perf_data = np.zeros((10, 4))
perf_eval_cb = PerfEvalCustomCallback(perf_data)
early_stop_cb = keras.callbacks.EarlyStopping(monitor='loss', mode='min', patience=4)

```

In [108...

```
model.summary()
```

Model: "model\_8"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_30 (InputLayer)	[(None, 100, 100, 3)]	0	[]
block1_conv1 (Conv2D)	(None, 49, 49, 32)	864	['input_30[0][0]']
block1_conv1_bn (BatchNormalization)	(None, 49, 49, 32)	128	['block1_conv1[0][0]']
block1_conv1_act (Activation)	(None, 49, 49, 32)	0	['block1_conv1_bn[0][0]']
block1_conv2 (Conv2D)	(None, 47, 47, 64)	18432	['block1_conv1_act[0][0]']
block1_conv2_bn (BatchNormalization)	(None, 47, 47, 64)	256	['block1_conv2[0][0]']
block1_conv2_act (Activation)	(None, 47, 47, 64)	0	['block1_conv2_bn[0][0]']
block2_sepconv1 (SeparableConv2D)	(None, 47, 47, 128)	8768	['block1_conv2_act[0][0]']
block2_sepconv1_bn (BatchNormalization)	(None, 47, 47, 128)	512	['block2_sepconv1[0][0]']
block2_sepconv2_act (Activation)	(None, 47, 47, 128)	0	['block2_sepconv1_bn[0][0]']
block2_sepconv2 (SeparableConv2D)	(None, 47, 47, 128)	17536	['block2_sepconv2_act[0][0]']
block2_sepconv2_bn (BatchNormalization)	(None, 47, 47, 128)	512	['block2_sepconv2[0][0]']
conv2d_16 (Conv2D)	(None, 24, 24, 128)	8192	['block1_conv2_act[0][0]']
block2_pool (MaxPooling2D)	(None, 24, 24, 128)	0	['block2_sepconv2_bn[0][0]']
batch_normalization_70 (BatchNormalization)	(None, 24, 24, 128)	512	['conv2d_16[0][0]']
add_48 (Add)	(None, 24, 24, 128)	0	['block2_pool[0][0]',

				'batch_normalization_7
0[0][0]']				
block3_sepconv1_act (Activation)	(None, 24, 24, 128)	0		['add_48[0][0]']
block3_sepconv1 (SeparableConv2D)	(None, 24, 24, 256)	33920		['block3_sepconv1_act[0][0]']
block3_sepconv1_bn (BatchNormalization)	(None, 24, 24, 256)	1024		['block3_sepconv1[0][0]']
block3_sepconv2_act (Activation)	(None, 24, 24, 256)	0		['block3_sepconv1_bn[0][0]']
block3_sepconv2 (SeparableConv2D)	(None, 24, 24, 256)	67840		['block3_sepconv2_act[0][0]']
block3_sepconv2_bn (BatchNormalization)	(None, 24, 24, 256)	1024		['block3_sepconv2[0][0]']
conv2d_17 (Conv2D)	(None, 12, 12, 256)	32768		['add_48[0][0]']
block3_pool (MaxPooling2D)	(None, 12, 12, 256)	0		['block3_sepconv2_bn[0][0]']
batch_normalization_71 (BatchNormalization)	(None, 12, 12, 256)	1024		['conv2d_17[0][0]']
add_49 (Add)	(None, 12, 12, 256)	0		['block3_pool[0][0]', 'batch_normalization_71[0][0]']
1[0][0]']				
block4_sepconv1_act (Activation)	(None, 12, 12, 256)	0		['add_49[0][0]']
block4_sepconv1 (SeparableConv2D)	(None, 12, 12, 728)	188672		['block4_sepconv1_act[0][0]']
block4_sepconv1_bn (BatchNormalization)	(None, 12, 12, 728)	2912		['block4_sepconv1[0][0]']
block4_sepconv2_act (Activation)	(None, 12, 12, 728)	0		['block4_sepconv1_bn[0][0]']
block4_sepconv2 (SeparableConv2D)	(None, 12, 12, 728)	536536		['block4_sepconv2_act[0][0]']
block4_sepconv2_bn (BatchNormalization)	(None, 12, 12, 728)	2912		['block4_sepconv2[0][0]']



conv2d_18 (Conv2D)	(None, 6, 6, 728)	186368	['add_49[0][0]']
block4_pool (MaxPooling2D)	(None, 6, 6, 728)	0	['block4_sepconv2_bn[0][0]']
batch_normalization_72 (Batch Normalization)	(None, 6, 6, 728)	2912	['conv2d_18[0][0]']
add_50 (Add)	(None, 6, 6, 728)	0	['block4_pool[0][0]', 'batch_normalization_72[0][0]']
block5_sepconv1_act (Activation)	(None, 6, 6, 728)	0	['add_50[0][0]']
block5_sepconv1 (SeparableConv2D)	(None, 6, 6, 728)	536536	['block5_sepconv1_act[0][0]']
block5_sepconv1_bn (Batch Normalization)	(None, 6, 6, 728)	2912	['block5_sepconv1[0][0]']
block5_sepconv2_act (Activation)	(None, 6, 6, 728)	0	['block5_sepconv1_bn[0][0]']
block5_sepconv2 (SeparableConv2D)	(None, 6, 6, 728)	536536	['block5_sepconv2_act[0][0]']
block5_sepconv2_bn (Batch Normalization)	(None, 6, 6, 728)	2912	['block5_sepconv2[0][0]']
block5_sepconv3_act (Activation)	(None, 6, 6, 728)	0	['block5_sepconv2_bn[0][0]']
block5_sepconv3 (SeparableConv2D)	(None, 6, 6, 728)	536536	['block5_sepconv3_act[0][0]']
block5_sepconv3_bn (Batch Normalization)	(None, 6, 6, 728)	2912	['block5_sepconv3[0][0]']
add_51 (Add)	(None, 6, 6, 728)	0	['block5_sepconv3_bn[0][0]', 'add_50[0][0]']
block6_sepconv1_act (Activation)	(None, 6, 6, 728)	0	['add_51[0][0]']
block6_sepconv1 (SeparableConv2D)	(None, 6, 6, 728)	536536	['block6_sepconv1_act[0][0]']
block6_sepconv1_bn (Batch Normalization)	(None, 6, 6, 728)	2912	['block6_sepconv1[0][0]']

```

[0]']
lization)

block6_sepconv2_act (Activatio (None, 6, 6, 728) 0 ['block6_sepconv1_bn[0]
[0]']
n)

block6_sepconv2 (SeparableConv (None, 6, 6, 728) 536536 ['block6_sepconv2_act
[0][0]']
2D)

block6_sepconv2_bn (BatchNorma (None, 6, 6, 728) 2912 ['block6_sepconv2[0]
[0]']
lization)

block6_sepconv3_act (Activatio (None, 6, 6, 728) 0 ['block6_sepconv2_bn[0]
[0]']
n)

block6_sepconv3 (SeparableConv (None, 6, 6, 728) 536536 ['block6_sepconv3_act
[0][0]']
2D)

block6_sepconv3_bn (BatchNorma (None, 6, 6, 728) 2912 ['block6_sepconv3[0]
[0]']
lization)

add_52 (Add) (None, 6, 6, 728) 0 ['block6_sepconv3_bn[0]
[0]',
'add_51[0][0]']

block7_sepconv1_act (Activatio (None, 6, 6, 728) 0 ['add_52[0][0]']
n)

block7_sepconv1 (SeparableConv (None, 6, 6, 728) 536536 ['block7_sepconv1_act
[0][0]']
2D)

block7_sepconv1_bn (BatchNorma (None, 6, 6, 728) 2912 ['block7_sepconv1[0]
[0]']
lization)

block7_sepconv2_act (Activatio (None, 6, 6, 728) 0 ['block7_sepconv1_bn[0]
[0]']
n)

block7_sepconv2 (SeparableConv (None, 6, 6, 728) 536536 ['block7_sepconv2_act
[0][0]']
2D)

block7_sepconv2_bn (BatchNorma (None, 6, 6, 728) 2912 ['block7_sepconv2[0]
[0]']
lization)

block7_sepconv3_act (Activatio (None, 6, 6, 728) 0 ['block7_sepconv2_bn[0]
[0]']
n)

block7_sepconv3 (SeparableConv (None, 6, 6, 728) 536536 ['block7_sepconv3_act
[0][0]']

```

2D)				
block7_sepconv3_bn (BatchNormalization)	(None, 6, 6, 728)	2912	['block7_sepconv3[0][0]']	
add_53 (Add)	(None, 6, 6, 728)	0	['block7_sepconv3_bn[0][0]', 'add_52[0][0]']	
block8_sepconv1_act (Activation)	(None, 6, 6, 728)	0	['add_53[0][0]']	
block8_sepconv1 (SeparableConv2D)	(None, 6, 6, 728)	536536	['block8_sepconv1_act[0][0]']	
block8_sepconv1_bn (BatchNormalization)	(None, 6, 6, 728)	2912	['block8_sepconv1[0][0]']	
block8_sepconv2_act (Activation)	(None, 6, 6, 728)	0	['block8_sepconv1_bn[0][0]']	
block8_sepconv2 (SeparableConv2D)	(None, 6, 6, 728)	536536	['block8_sepconv2_act[0][0]']	
block8_sepconv2_bn (BatchNormalization)	(None, 6, 6, 728)	2912	['block8_sepconv2[0][0]']	
block8_sepconv3_act (Activation)	(None, 6, 6, 728)	0	['block8_sepconv2_bn[0][0]']	
block8_sepconv3 (SeparableConv2D)	(None, 6, 6, 728)	536536	['block8_sepconv3_act[0][0]']	
block8_sepconv3_bn (BatchNormalization)	(None, 6, 6, 728)	2912	['block8_sepconv3[0][0]']	
add_54 (Add)	(None, 6, 6, 728)	0	['block8_sepconv3_bn[0][0]', 'add_53[0][0]']	
block9_sepconv1_act (Activation)	(None, 6, 6, 728)	0	['add_54[0][0]']	
block9_sepconv1 (SeparableConv2D)	(None, 6, 6, 728)	536536	['block9_sepconv1_act[0][0]']	
block9_sepconv1_bn (BatchNormalization)	(None, 6, 6, 728)	2912	['block9_sepconv1[0][0]']	

block9_sepconv2_act (Activation)	(None, 6, 6, 728)	0	['block9_sepconv1_bn[0]
block9_sepconv2 (SeparableConv2D)	(None, 6, 6, 728)	536536	['block9_sepconv2_act[0][0]']
block9_sepconv2_bn (BatchNormalization)	(None, 6, 6, 728)	2912	['block9_sepconv2[0]
block9_sepconv3_act (Activation)	(None, 6, 6, 728)	0	['block9_sepconv2_bn[0]
block9_sepconv3 (SeparableConv2D)	(None, 6, 6, 728)	536536	['block9_sepconv3_act[0][0]']
block9_sepconv3_bn (BatchNormalization)	(None, 6, 6, 728)	2912	['block9_sepconv3[0]
add_55 (Add)	(None, 6, 6, 728)	0	['block9_sepconv3_bn[0]
			'add_54[0][0]']
block10_sepconv1_act (Activation)	(None, 6, 6, 728)	0	['add_55[0][0]']
block10_sepconv1 (SeparableConv2D)	(None, 6, 6, 728)	536536	['block10_sepconv1_act[0][0]']
block10_sepconv1_bn (BatchNormalization)	(None, 6, 6, 728)	2912	['block10_sepconv1[0]
block10_sepconv2_act (Activation)	(None, 6, 6, 728)	0	['block10_sepconv1_bn[0][0]']
block10_sepconv2 (SeparableConv2D)	(None, 6, 6, 728)	536536	['block10_sepconv2_act[0][0]']
block10_sepconv2_bn (BatchNormalization)	(None, 6, 6, 728)	2912	['block10_sepconv2[0]
block10_sepconv3_act (Activation)	(None, 6, 6, 728)	0	['block10_sepconv2_bn[0][0]']
block10_sepconv3 (SeparableConv2D)	(None, 6, 6, 728)	536536	['block10_sepconv3_act[0][0]']
block10_sepconv3_bn (BatchNormalization)	(None, 6, 6, 728)	2912	['block10_sepconv3[0]

```

[0]']
alization)

add_56 (Add) (None, 6, 6, 728) 0 ['block10_sepconv3_bn
[0][0]', 'add_55[0][0]']

block11_sepconv1_act (Activati (None, 6, 6, 728) 0 ['add_56[0][0]']
on)

block11_sepconv1 (SeparableCon (None, 6, 6, 728) 536536 ['block11_sepconv1_act
[0][0]']
v2D)

block11_sepconv1_bn (BatchNorm (None, 6, 6, 728) 2912 ['block11_sepconv1[0]
[0]']
alization)

block11_sepconv2_act (Activati (None, 6, 6, 728) 0 ['block11_sepconv1_bn
[0][0]']
on)

block11_sepconv2 (SeparableCon (None, 6, 6, 728) 536536 ['block11_sepconv2_act
[0][0]']
v2D)

block11_sepconv2_bn (BatchNorm (None, 6, 6, 728) 2912 ['block11_sepconv2[0]
[0]']
alization)

block11_sepconv3_act (Activati (None, 6, 6, 728) 0 ['block11_sepconv2_bn
[0][0]']
on)

block11_sepconv3 (SeparableCon (None, 6, 6, 728) 536536 ['block11_sepconv3_act
[0][0]']
v2D)

block11_sepconv3_bn (BatchNorm (None, 6, 6, 728) 2912 ['block11_sepconv3[0]
[0]']
alization)

add_57 (Add) (None, 6, 6, 728) 0 ['block11_sepconv3_bn
[0][0]', 'add_56[0][0]']

block12_sepconv1_act (Activati (None, 6, 6, 728) 0 ['add_57[0][0]']
on)

block12_sepconv1 (SeparableCon (None, 6, 6, 728) 536536 ['block12_sepconv1_act
[0][0]']
v2D)

block12_sepconv1_bn (BatchNorm (None, 6, 6, 728) 2912 ['block12_sepconv1[0]
[0]']
alization)

block12_sepconv2_act (Activati (None, 6, 6, 728) 0 ['block12_sepconv1_bn
[0][0]']
on)

```

block12_sepconv2 (SeparableConv2D)	(None, 6, 6, 728)	536536	['block12_sepconv2_activation[0][0]']
block12_sepconv2_bn (BatchNormalization)	(None, 6, 6, 728)	2912	['block12_sepconv2[0][0]']
block12_sepconv3_act (Activation)	(None, 6, 6, 728)	0	['block12_sepconv2_bn[0][0]']
block12_sepconv3 (SeparableConv2D)	(None, 6, 6, 728)	536536	['block12_sepconv3_activation[0][0]']
block12_sepconv3_bn (BatchNormalization)	(None, 6, 6, 728)	2912	['block12_sepconv3[0][0]']
add_58 (Add)	(None, 6, 6, 728)	0	['block12_sepconv3_bn[0][0]', 'add_57[0][0]']
block13_sepconv1_act (Activation)	(None, 6, 6, 728)	0	['add_58[0][0]']
block13_sepconv1 (SeparableConv2D)	(None, 6, 6, 728)	536536	['block13_sepconv1_activation[0][0]']
block13_sepconv1_bn (BatchNormalization)	(None, 6, 6, 728)	2912	['block13_sepconv1[0][0]']
block13_sepconv2_act (Activation)	(None, 6, 6, 728)	0	['block13_sepconv1_bn[0][0]']
block13_sepconv2 (SeparableConv2D)	(None, 6, 6, 1024)	752024	['block13_sepconv2_activation[0][0]']
block13_sepconv2_bn (BatchNormalization)	(None, 6, 6, 1024)	4096	['block13_sepconv2[0][0]']
conv2d_19 (Conv2D)	(None, 3, 3, 1024)	745472	['add_58[0][0]']
block13_pool (MaxPooling2D)	(None, 3, 3, 1024)	0	['block13_sepconv2_bn[0][0]']
batch_normalization_73 (BatchNormalization)	(None, 3, 3, 1024)	4096	['conv2d_19[0][0]']
add_59 (Add)	(None, 3, 3, 1024)	0	['block13_pool[0][0]', 'batch_normalization_73[0][0]']

block14_sepconv1 (SeparableConv2D)	(None, 3, 3, 1536)	1582080	['add_59[0][0]']
block14_sepconv1_bn (BatchNormalization)	(None, 3, 3, 1536)	6144	['block14_sepconv1[0][0]']
block14_sepconv1_act (Activation)	(None, 3, 3, 1536)	0	['block14_sepconv1_bn[0][0]']
block14_sepconv2 (SeparableConv2D)	(None, 3, 3, 2048)	3159552	['block14_sepconv1_act[0][0]']
block14_sepconv2_bn (BatchNormalization)	(None, 3, 3, 2048)	8192	['block14_sepconv2[0][0]']
block14_sepconv2_act (Activation)	(None, 3, 3, 2048)	0	['block14_sepconv2_bn[0][0]']
global_average_pooling2d_9 (GlobalAveragePooling2D)	(None, 2048)	0	['block14_sepconv2_act[0][0]']
dense_20 (Dense)	(None, 200)	409800	['global_average_pooling2d_9[0][0]']
dense_21 (Dense)	(None, 100)	20100	['dense_20[0][0]']
dense_22 (Dense)	(None, 50)	5050	['dense_21[0][0]']
dense_23 (Dense)	(None, 10)	510	['dense_22[0][0]']

```

=====
Total params: 21,296,940
Trainable params: 21,242,412
Non-trainable params: 54,528

```



In [106...]

```
hobj = model.fit(x_train, labels_train_one, validation_data=(x_val, labels_val_one), epochs=10, shuffle=True, callbacks=[perf_eval_cb, early_stop_cb], verbose=1)
```

Epoch 1/10

```
377/377 [=====] - 791s 2s/step - loss: 0.9091 - accuracy: 0.7068 - val_loss: 1.6019 - val_accuracy: 0.6823
```

Epoch 2/10

```
377/377 [=====] - 765s 2s/step - loss: 0.3833 - accuracy: 0.8868 - val_loss: 0.6148 - val_accuracy: 0.8199
```

Epoch 3/10

```
377/377 [=====] - 778s 2s/step - loss: 0.2659 - accuracy: 0.9196 - val_loss: 0.7714 - val_accuracy: 0.8643
```

Epoch 4/10

```

377/377 [=====] - 783s 2s/step - loss: 0.2030 - accuracy: 0.939
5 - val_loss: 0.7210 - val_accuracy: 0.8584
Epoch 5/10
377/377 [=====] - 811s 2s/step - loss: 0.1621 - accuracy: 0.950
4 - val_loss: 0.4860 - val_accuracy: 0.8663
Epoch 6/10
377/377 [=====] - 799s 2s/step - loss: 0.1367 - accuracy: 0.958
1 - val_loss: 0.3396 - val_accuracy: 0.9008
Epoch 7/10
377/377 [=====] - 785s 2s/step - loss: 0.1059 - accuracy: 0.968
8 - val_loss: 0.3019 - val_accuracy: 0.9194
Epoch 8/10
377/377 [=====] - 751s 2s/step - loss: 0.1004 - accuracy: 0.969
8 - val_loss: 0.2482 - val_accuracy: 0.9271
Epoch 9/10
377/377 [=====] - 763s 2s/step - loss: 0.0850 - accuracy: 0.974
9 - val_loss: 0.2475 - val_accuracy: 0.9392
Epoch 10/10
377/377 [=====] - 756s 2s/step - loss: 0.0812 - accuracy: 0.976
1 - val_loss: 0.2512 - val_accuracy: 0.9382

```

```

-----
NameError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_11052\2076790798.py in <module>
      2         shuffle=True, callbacks=[perf_eval_cb, early_stop_cb], verbose=1)
      3     eff_epochs = len(hobj.history['loss'])
----> 4     eval_data = (perf_data[0:eff_epochs,:], dataset)

```

**NameError:** name 'dataset' is not defined

```

In [106...   eff_epochs = len(hobj.history['loss'])
           eval_data = (perf_data[0:eff_epochs,:], cnn_dataset)

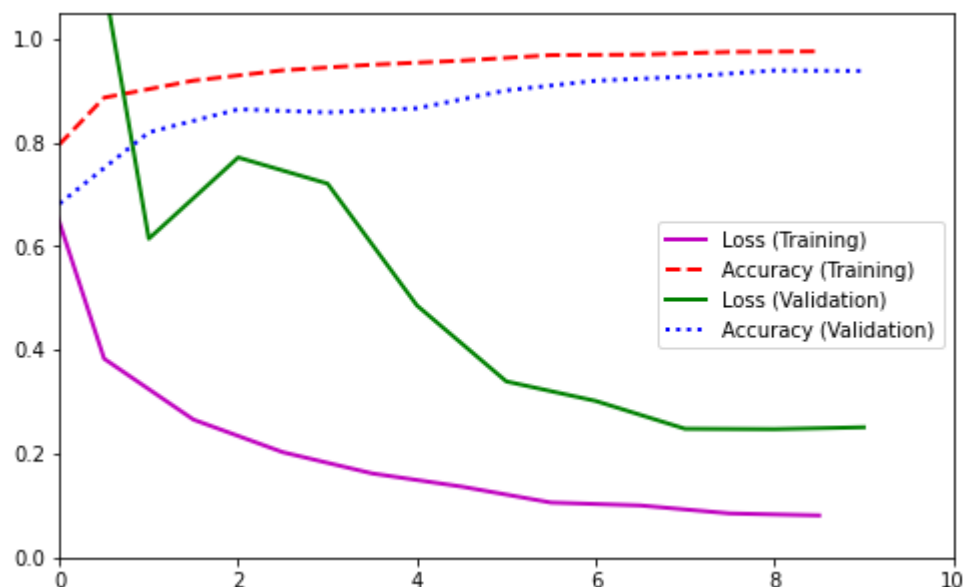
```

```

In [106...   evaluate_model("test", model, eval_data)

```

[test] Training Accuracy: 97.784%, Validation Accuracy: 93.824%



```

126/126 [=====] - 42s 242ms/step - loss: 0.2857 - accuracy: 0.9
335
[test] Test loss: 0.28567; test accuracy: 93.353%
[test] Test cross entropy loss: 0.28447

```



# FULL TRAINING

In [107... `cnn_input.shape`

Out[107... (26880, 100, 100, 3)

In [107... `labels_new_ohe = OneHotEncoder().fit_transform(labels_new.reshape(-1,1)).toarray()`

In [107... `labels_new_ohe.shape`

Out[107... (26880, 10)

In [109... `hobj = model.fit(cnn_input, labels_new_ohe, epochs=8, batch_size=100, shuffle=True, verbose=1)`

```
Epoch 1/8
269/269 [=====] - 967s 4s/step - loss: 0.2659 - accuracy: 0.9217
Epoch 2/8
269/269 [=====] - 960s 4s/step - loss: 0.1679 - accuracy: 0.9480
Epoch 3/8
269/269 [=====] - 975s 4s/step - loss: 0.1285 - accuracy: 0.9607
Epoch 4/8
269/269 [=====] - 964s 4s/step - loss: 0.0978 - accuracy: 0.9718
Epoch 5/8
269/269 [=====] - 1004s 4s/step - loss: 0.0806 - accuracy: 0.9766
Epoch 6/8
1/269 [.....] - ETA: 16:14 - loss: 0.0712 - accuracy: 0.9600
```

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_11052\1563957083.py in <module>
----> 1 hobj = model.fit(cnn_input, labels_new_ohe, epochs=8, batch_size=100,
      2                 shuffle=True, verbose=1)

~\anaconda3\lib\site-packages\keras\utils\traceback_utils.py in error_handler(*args, **k
wargs)
    62     filtered_tb = None
    63     try:
--> 64         return fn(*args, **kwargs)
    65     except Exception as e: # pylint: disable=broad-except
    66         filtered_tb = _process_traceback_frames(e.__traceback__)

~\anaconda3\lib\site-packages\keras\engine\training.py in fit(self, x, y, batch_size, ep
ochs, verbose, callbacks, validation_split, validation_data, shuffle, class_weight, samp
le_weight, initial_epoch, steps_per_epoch, validation_steps, validation_batch_size, vali
dation_freq, max_queue_size, workers, use_multiprocessing)
   1382         _r=1):
   1383             callbacks.on_train_batch_begin(step)
```

```

-> 1384         tmp_logs = self.train_function(iterator)
1385         if data_handler.should_sync:
1386             context.async_wait()

~\anaconda3\lib\site-packages\tensorflow\python\util\traceback_utils.py in error_handler
(*args, **kwargs)
148     filtered_tb = None
149     try:
-> 150         return fn(*args, **kwargs)
151     except Exception as e:
152         filtered_tb = _process_traceback_frames(e.__traceback__)

~\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py in __call__(self,
*args, **kws)
913
914     with OptionalXlaContext(self._jit_compile):
-> 915         result = self._call(*args, **kws)
916
917         new_tracing_count = self.experimental_get_tracing_count()

~\anaconda3\lib\site-packages\tensorflow\python\eager\def_function.py in _call(self, *ar
gs, **kws)
945         # In this case we have created variables on the first call, so we run the
946         # defunned version which is guaranteed to never create variables.
-> 947         return self._stateless_fn(*args, **kws) # pylint: disable=not-callable
948     elif self._stateful_fn is not None:
949         # Release the lock early so that multiple threads can perform the call

~\anaconda3\lib\site-packages\tensorflow\python\eager\function.py in __call__(self, *arg
s, **kwargs)
2954         (graph_function,
2955          filtered_flat_args) = self._maybe_define_function(args, kwargs)
-> 2956         return graph_function._call_flat(
2957             filtered_flat_args, captured_inputs=graph_function.captured_inputs) # p
ylint: disable=protected-access
2958

~\anaconda3\lib\site-packages\tensorflow\python\eager\function.py in _call_flat(self, ar
gs, captured_inputs, cancellation_manager)
1851         and executing_eagerly):
1852         # No tape is watching; skip to running the function.
-> 1853         return self._build_call_outputs(self._inference_function.call(
1854             ctx, args, cancellation_manager=cancellation_manager))
1855         forward_backward = self._select_forward_and_backward_functions(

~\anaconda3\lib\site-packages\tensorflow\python\eager\function.py in call(self, ctx, arg
s, cancellation_manager)
497         with _InterpolateFunctionError(self):
498             if cancellation_manager is None:
-> 499                 outputs = execute.execute(
500                     str(self.signature.name),
501                     num_outputs=self._num_outputs,

~\anaconda3\lib\site-packages\tensorflow\python\eager\execute.py in quick_execute(op_nam
e, num_outputs, inputs, attrs, ctx, name)
52     try:
53         ctx.ensure_initialized()
---> 54         tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name, op_name,
55             inputs, attrs, num_outputs)
56     except core._NotOkStatusException as e:

```

**KeyboardInterrupt:**

In [109...]

```
model.save('final_trained_model')
```

INFO:tensorflow:Assets written to: final\_trained\_model\assets

## TEST SECTION

In [ ]:

```
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
import tensorflow.keras as keras
import cv2
from sklearn.preprocessing import MinMaxScaler, OneHotEncoder
from sklearn.model_selection import train_test_split
import time
```

In [ ]:

```
test_data = np.load("put_path_here.npy")
test_labels = np.load("put_path_here.npy")

weights = "final_trained_model"
model = keras.models.load_model(weights)
```

## Functions used to preprocess data

In [ ]:

```
def show(im):
    plt.imshow(im, cmap='gray')
    plt.show()

def invert(im):
    return (im*-1)+255

def min_max_scale(im):
    im = MinMaxScaler().fit_transform(im.ravel().reshape(-1,1))
    return im.reshape(300,300)

def brighten(im):
    m = np.max(im)-.3
    im[im >= m] += 0.3
    im[im < m] -= 0.3
    return im

def blur(im, kernel):
    im = im.astype("uint8")
    im = cv2.medianBlur(im, kernel)
    return im

def morph_close(im, kernel, i):
    return cv2.morphologyEx(im, cv2.MORPH_CLOSE, kernel, iterations=i)

def morph_open(im, kernel, i):
    return cv2.morphologyEx(im, cv2.MORPH_OPEN, kernel, iterations=i)
```

```

def morph_dilate(im, kernel, i):
    return cv2.dilate(im, kernel, iterations=i)

def morph_erode(im, kernel, i):
    return cv2.erode(im, kernel, iterations=i)

def transform(im):
    dest = np.float32([[0,0],[300,0],[0,300],[300,300]])
    source = np.float32([[25,25],[275,25],[25,275],[275,275]])
    res = cv2.getPerspectiveTransform(source,dest)
    return cv2.warpPerspective(im, res, (300,300))

def preprocess(ims):
    r = []
    for im in ims.T:
        im = im.reshape((300,300))
        im = transform(im)
        im = invert(im)
        im = blur(im, 5)
        im = morph_erode(im, (3,3), 5)
        im = morph_dilate(im, (3,3), 5)
        im = morph_close(im, (3,3), 5)
        im = morph_open(im, (3,3), 5)
        im = min_max_scale(im)
        im = brighten(im)
        im = min_max_scale(im)
        im = cv2.resize(im, (100,100))
        im = np.stack((im,)*3, axis=-1)
        r.append(im)
    return np.array(r)

```

```

In [ ]: test_data = preprocess(test_data)
        test_labels_ohe = OneHotEncoder().fit_transform(test_labels.reshape(-1,1)).toarray()

```

```

In [ ]: model = keras.models.load_model("put_model_file_here")

```

```

In [ ]: test_preds = model.predict(test_data)

        accuracy_obj = keras.metrics.CategoricalAccuracy()
        accuracy_obj.update_state(test_labels_ohe, test_preds)
        test_acc = accuracy_obj.result()

        test_loss, _ = model.evaluate(test_data, test_labels_ohe)

        cat_loss_obj = tf.keras.losses.CategoricalCrossentropy()
        test_ce_loss = cat_loss_obj(test_labels_ohe, test_preds)

        print('[{}] Test loss: {:.5f}; test accuracy: {:.3f}%'.format(name, test_loss, 100*tes
        print('[{}] Test cross entropy loss: {:.5f}'.format(name, test_ce_loss))

```